

What one sweep costs on a commercial API

Every dataset in the suite at 50 records, priced against the cheap tier of GPT, Claude and Gemini — from token counts the servers actually reported, not from guesses.

\$19

GPT-5 NANO

\$22

GEMINI 2.5 FLASH-LITE

\$260

CLAUDE HAIKU 4.5

62.7M

INPUT TOKENS

39.4M

OUTPUT TOKENS

Fifty records is not fifty calls

A dataset does not cost one request per record. Each one is asked in both elicitations, some pose two tasks, every question is repeated three times for self-consistency, and the interactive benchmarks turn a single record into a whole conversation.

44	94	14,100	29,679	6,600
DATASETS	TASKS	ANSWERS SCORED	MODEL CALLS	JUDGE CALLS
every benchmark in the suite	io / cot × the tasks each poses	tasks × 50 records × 3 repeats	episodes expand into turns	22 datasets have no checkable answer

Two elicitations

io answers straight away; cot reasons first. They are separate tasks with separate scores — and wildly different output bills.

Three repeats

Self-consistency and Best-of-N are a plurality over k samples. At k=1 they do not exist, so every record is asked three times at temperature 0.7.

Five interactive benchmarks

cloud_opsbench, ddxplus, med_inquire, medqdx, vivabench run 9–14 turns per record, each turn re-sending the whole transcript. They are 5 of 44 datasets and **82% of the input bill**.

Measured, then stated with its error bars

Where the counts come from

Every figure is the `usage.prompt_tokens / usage.completion_tokens` the serving engine returned, summed over each task's raw payloads and divided by its records. Nothing is estimated with a client-side tokenizer.

Tokenizers differ

Counts are in Qwen's tokenizer. GPT, Claude and Gemini each segment text slightly differently; on English prose the spread is within roughly $\pm 10\%$, which is smaller than the output-length uncertainty above.

Coverage

33 datasets come from the 50-record $\times$ 3-repeat run of 14 Sep 2026. The remaining 11 were measured the same way on a 3-record probe the same day — per-record means, so they scale identically.

What will move

Output length is a property of the model, not the dataset. These lengths are Qwen3.5-2B's. A different model reasons at a different length, so treat output as the soft half of the estimate — the sensitivity slide prices that.

Average tokens per record, dataset by dataset

One row per dataset. **In/record** covers every API call a record costs, so for an interactive benchmark it is the whole episode. Output is split by elicitation, because that is where the two differ.

DATASET	TASKS	CALLS/REC	IN/REC	OUT/REC	· IO	OUT/REC	· COT	TOTAL IN	TOTAL OUT	GPT-5 NANO	HAIKU 4.5
ddxplus INT	2	9.4	39,302	15,703	CAP	16,814	CAP	11.8M	4.9M	\$2.54	\$36.18
abd	2	1	4,600	8,862	CAP	23,228	CAP	1.4M	4.8M	\$1.99	\$25.45
proof_writer 3-REC	2	1	385	10,672	CAP	21,677	CAP	0.1M	4.9M	\$1.95	\$24.38
scir 3-REC	2	1	4,820		4	22,644	CAP	1.4M	3.4M	\$1.43	\$18.43
neurL JUDGED 3-REC	2	1	613		16	21,595	CAP	0.2M	3.2M	\$1.32	\$16.61
vivabench INT JUDGED 3-REC	2	13.5	60,456		3,634		3,074	18.1M	1.0M	\$1.32	\$23.39
true_detective 3-REC	2	1	1,636		4	21,335	CAP	0.5M	3.2M	\$1.30	\$16.49
uniadilr_hgc 3-REC	2	1	1,030		37	21,372	CAP	0.3M	3.2M	\$1.30	\$16.37
synpat 3-REC	2	1	544		38	21,340	CAP	0.2M	3.2M	\$1.29	\$16.20
hypobench JUDGED	2	1	1,474		79	8,152	CAP	0.4M	1.2M	\$0.53	\$6.83
house_md JUDGED	2	1	593		460	7,460	CAP	0.2M	1.2M	\$0.50	\$6.34
cloud_opsbench INT JUDGED	2	14	26,766		289		277	8.0M	0.1M	\$0.45	\$8.67
aer	2	1	888		222	6,706	CAP	0.3M	1.0M	\$0.43	\$5.46

DATASET	TASKS	CALLS/REC	IN/REC	OUT/REC · IO	OUT/REC · COT	TOTAL IN	TOTAL OUT	GPT-5 NANO	HAIKU 4.5
musr	2	1	1,348	5	212	0.4M	0.0M	\$0.03	\$0.57
aiops2025	2	1	1,117	8	126	0.3M	0.0M	\$0.02	\$0.44
b_copa	2	1	188	5	360	0.1M	0.1M	\$0.02	\$0.33
climate_fever	2	1	480	4	200	0.1M	0.0M	\$0.02	\$0.30
commonwhy JUDGED	2	1	208	41	95	0.1M	0.0M	\$0.02	\$0.38
copa	2	1	186	5	324	0.1M	0.0M	\$0.02	\$0.30
uncommonsense JUDGED 3-REC	2	1	259	63	76	0.1M	0.0M	\$0.02	\$0.40
xcopa 3-REC	2	1	216	5	4	0.1M	0.0M	\$0.00	\$0.07
44 datasets	94		4,145			58.4M	39.3M	\$18.91	\$259.96

INTERACTIVE multi-turn episode SEQUENTIAL staged delivery JUDGED needs an LLM grader CAP ran to the token ceiling

3-REC measured on the 3-record probe · totals are for 50 records × 3 repeats × the dataset's tasks

20 MEDIAN OUT/REC · IO the answer line, and little else	716 MEDIAN OUT/REC · COT reasoning, then the answer line	×5 MORE OUTPUT OVERALL across every dataset that runs both	22 EXTRA IN/REC · COT the instruction itself is cheap
---	--	--	---

scir io 4 → cot 22,644		+22,640
neulr io 16 → cot 21,595		+21,579
uniadilr_hgc io 37 → cot 21,372		+21,335
true_detective io 4 → cot 21,335		+21,331
synpat io 38 → cot 21,340		+21,302
abd io 8,862 → cot 23,228		+14,366
proof_writer io 10,672 → cot 21,677		+11,005
hypobench io 70 → cot 8,452		+8,073

What the whole suite weighs


58.4M INPUT TOKENS prompts, every turn of every episode	39.3M OUTPUT TOKENS completions, reasoning included	4.3M JUDGE INPUT bounded: 600-char candidate, 200-word case	1.5:1 INPUT : OUTPUT decides which of the two prices matters
--	--	--	---

The judge is a rounding error

6,600 judge calls add 4.3M input and 99,000 output tokens — \$0.25 on GPT-5 nano, \$4.79 on Haiku 4.5. The adapters clip the candidate to 600 characters and the case to 200 words before the judge sees them, so this line cannot run away. Grading with a different model from the one under test stays mandatory either way.


One full sweep, priced five ways

List prices as of September 2026, cheap tier per vendor. Solid bar is the sweep; the faded extension is the LLM judge on the 21 datasets that need one.



MODEL	\$/1M IN	\$/1M OUT	SWEEP	JUDGE	TOTAL	BATCH API -50%
GPT-5 nano · OpenAI	\$0.05	\$0.40	\$18.66	\$0.25	\$18.91	\$9.46

Five datasets, 82% of the input



01	vivabench alone is 31% of all input — 13.5 calls per record, each re-sending the transcript.	\$1.32
02	The four interactive benchmarks are 82% of input from 11% of the datasets.	\$4.92
03	ddxplus and proof_writer are 25% of all output — long formal answers under chain-of-thought.	\$4.49
04	The median dataset costs well under a dollar on the cheap tier. Most of the suite is free in practice.	\$0.11

Three ways to pay less for the same run

- 01 **Batch API** — all three vendors halve input and output for asynchronous work. The engine already sends grouped batch requests; the change is the endpoint and a 24-hour turnaround. **–50%**
- 02 **Prompt caching on the episodes** — 79–88% of every interactive turn is transcript that was already sent, and a cache read bills at a tenth. OpenAI caches automatically with no write charge; Anthropic bills a 1.25× write the first time a prefix is seen. **–10%**
- 03 **Drop the third repeat** — $k=3$ is the minimum that gives self-consistency and Best-of-N a meaning. $k=2$ keeps a vote, at two-thirds of the bill; $k=1$ loses both metrics entirely. **–33%**

The output half of this estimate is the soft half

Input is fixed by the prompts. Output is not — it is whatever the model chooses to write, and Qwen3.5-2B wrote until it hit its ceiling on 11 datasets under chain-of-thought. **Cap output at 4,096 tokens** and the same sweep costs \$8.70 on GPT-5 nano and \$132.30 on Haiku 4.5, against \$18.91 and \$259.96 as measured. Double the measured lengths instead and they become \$34.69 and \$457.18. The ranking between vendors never changes.

Run all three. It costs less than one day of the GPU.

\$19 GPT-5 NANO the whole suite, judge included	\$22 GEMINI 2.5 FLASH-LITE same suite, second vendor	\$260 CLAUDE HAIKU 4.5 same suite, third vendor	\$301 ALL THREE one comparable grid, three families
--	---	--	--

Start with GPT-5 nano

At \$18.91 for the full 44-dataset sweep it is cheap enough to run first and discover the operational problems — rate limits, refusals, truncation — before spending real money.

Add Gemini 2.5 Flash-Lite before 16 Oct

At \$22.05 it is within rounding distance of nano and retires on 16 October 2026. After that the cheapest Gemini is 3.1 Flash-Lite at \$74.85.

Budget Haiku 4.5 deliberately

Haiku 4.5 is 14× nano at \$259.96 — it prices input at \$1.00/M against nano’s \$0.05, on a suite that runs 1.5:1 input to output. Batch halves it to \$129.98, and a 4,096-token output ceiling brings it to \$132.30.

Token counts measured 14 Sep 2026 from AbductionBench run 20260914-185857_full and probe costprobe, on Qwen3.5-2B served by vLLM. Prices are vendor list prices for September 2026 and exclude prompt caching unless stated. Sources: [OpenAI](#), [Anthropic](#), [Google](#).